



# MongoDB

---

## An Introduction



# Who's Waldo

- Geek
- Problem solver
- Systems Engineer
- @gwaldo
- [github.com/gwaldo](https://github.com/gwaldo)





# Caveats

- No affiliation with I0gen.
- Language



# What is MongoDB?

- “MongoDB (from "humongous") is a scalable, high-performance, open source NoSQL database.”





# Features! Quick!

- JSON Document-oriented
  - (actually BSON)



# Features! Quick!

- Indexing





# Features! Quick!

- Querying

```
db.geeks.find({name:"waldo"}, {sexy: 1})
```



# Features! Quick!

- Replication





# Features! Quick!

- Sharding



# More Features

- Journaling





# More Features

- **Write Concern**



# More Features

- 'Stored Procedures'

```
function addNumbers ( x , y ) {  
    return x + y;  
}
```

```
db.system.js.save({_id:"addNumbers",  
value:function(x, y){ return x + y; }});
```





# More Features

- 'Stored Procedures'

```
> db.eval( 'addNumbers(17, 25)' );  
42
```



# More Features







# More Features

- Seriously, ...



# MongoDB is a \_\_\_\_\_ Database

- Document
- Open-Source
- High-Performance
- Horizontally-Scalable
- Full-Featured





# MongoDB is a Document Database

- Not .pdf or .doc files
- JSON objects
- Associative Arrays (PHP array, Python Dict, Ruby & Perl Hash, etc)



# MongoDB is an Open-Source Database

- <https://github.com/mongodb>
- AGPL license
- Originated and Sponsored by 10gen
- Commercial licenses available
- Contributions welcome





# MongoDB is a High-Performance Database

- C++
- Extensive use of MMapped-files
- Runs in ALL THE PLACES
- Data serialized as BSON



# MongoDB is a Horizontally-Scalable Database

- Replication
- Sharding
- Both Dead-Simple





# MongoDB is a Full-Featured Database

- Rich Querying
- Real-time Aggregation
- Traditionally Consistent
- Geospatial



# “NoSQL?”

- Non-Relational
- Flexible (if any) Schema
- not-ACID-ic
- Does not use SQL
  - uses a JSON Query style





# Platforms

- 64- and 32-bit
  - Don't use 32-bit
- \*nix, Mac, & Windows
- Binary, source, and package managers



# Officially Supported Languages

- MongoDB Supported:
  - C & C++
  - Erlang
  - Haskell
  - Java
  - JavaScript
- .NET
- Perl
- PHP
- Python
- Ruby
- Scala





# Community-Supported Languages

- Too many for me to list
- [http://www.mongodb.org/display/DOCS/  
Drivers](http://www.mongodb.org/display/DOCS/Drivers)



# SQL to MongoDB

| MySQL term  | Mongo term/concept    |
|-------------|-----------------------|
| database    | database              |
| table       | collection            |
| index       | index                 |
| row         | BSON document         |
| column      | BSON field            |
| join        | embedding and linking |
| primary key | _id field             |
| group by    | aggregation           |





| SQL Statement   | Mongo Statement  |
|---|--|
| <pre>CREATE TABLE USERS (a Number, b Number)</pre>        | implicit; can also be done explicitly with<br><pre>db.createCollection("mycoll")</pre> |
| <pre>ALTER TABLE users ADD ...</pre>                      | implicit   |
|   |  |
| <pre>INSERT INTO USERS VALUES (3,5)</pre>                 | <pre>db.users.insert({a:3,b:5})</pre>  |
|   |  |
| <pre>SELECT a,b FROM users</pre>                          | <pre>db.users.find({}, {a:1,b:1})</pre>  |
| <pre>SELECT * FROM users</pre>                            | <pre>db.users.find()</pre>   |
| <pre>SELECT * FROM users WHERE age=33</pre>               | <pre>db.users.find({age:33})</pre>   |
| <pre>SELECT a,b FROM users WHERE age=33</pre>             | <pre>db.users.find({age:33}, {a:1,b:1})</pre>  |
| <pre>SELECT * FROM users WHERE age=33 ORDER BY name</pre> | <pre>db.users.find({age:33}).sort({name:1})</pre>                                      |



# What it's Good at

- Archiving & Event Logging
- Documents / Content Management
- Gaming
- Mobile & Location Services
- Agile Development
- Real-Time Stats / Analysis





# What it's Bad at

- Complex Transactional (Banking & Accounting)
- Traditional Data Warehousing
- Where you absolutely need SQL (complex joins)



# On Joins

- Data Design
  - by-ref
  - by copy
- Separate queries in app logic





# Caveats & Gotchas

- Global Write Lock
  - On Writes, read first



# Caveats & Gotchas

- Queries are case-sensitive
  - `var test1 = db.test.find({'tags': 'jquery'}).count();`
  - `var test2 = db.test.find({'tags': 'jQuery'}).count();`
  - `test1 == test2; // Output is false - they do not query for the same information`





# Caveats & Gotchas

- Don't store numbers as strings
  - `{'count': 102};` // 'count' is stored as an int
  - `{'count': "102"};` // 'count' is stored as a string



# Caveats & Gotchas

- Document sizes capped at 16MB
  - Not a problem for much of the world...
  - but for the rest, GridFS





# Other tips

- Unless speed is paramount,
  - `getLastError`
- Use `.limit()` when using `.find()`
- When doing mass-updates, narrow the search AMAP



# Cool Tools

- ``mongostat``
- `db.serverStatus()`
- `db.stats()` & `db.<collection_name>.stats()`
- `db.printReplicationInfo()`
- `db.printSlaveReplicationInfo`
- `<query>.explain()`





# What to watch for

- Page Faults
- Index Misses
- Queue Length



# Massive Cop-Out







`_id`

- **primary key**
- **Automatically created**
- **Automatically indexed**
- **you may override if desired**



# ObjectID

- special 12-byte value
- Guaranteed unique across cluster





# Shell vs Drivers

- Don't build your apps in the shell



# Officially Supported Languages

- MongoDB Supported:
  - C & C++
  - Erlang
  - Haskell
  - Java
  - JavaScript
- .NET
- Perl
- PHP
- Python
- Ruby
- Scala





# Now what?

- Lots of concepts introduced



# Now what?

- Documentation is awesome
- [l0gen.com/presentations](http://l0gen.com/presentations)
- “Snail in a Turtleneck” blog
  - Kristina Chodorow





# Now what?

- Play!
- Talk about it!